

```
#Checking if GPU is running or not
```

```
!nvidia-smi
```

```
📄 Fri Sep 1 05:24:46 2023
+-----+
| NVIDIA-SMI 525.105.17   Driver Version: 525.105.17   CUDA Version: 12.0   |
+-----+-----+-----+-----+-----+
| GPU  Name      Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M. |
+-----+-----+-----+-----+-----+
|   0   Tesla T4            Off | 00000000:00:04:0  Off |           0          |
| N/A   75C    P8      14W / 70W |  0MiB / 15360MiB |           0%      Default |
+-----+-----+-----+-----+-----+

+-----+
| Processes: |
| GPU  GI  CI           PID  Type  Process name                        GPU Memory |
|          ID  ID                                   Usage   |
+-----+
| No running processes found |
+-----+
```

```
!pip install datasets transformers[sentencepiece] sacrebleu -q
```

```
# Importing the Required Libraries
import os
import sys
import transformers
import tensorflow as tf
from datasets import load_dataset
from transformers import AutoTokenizer
from transformers import TFAutoModelForSeq2SeqLM, DataCollatorForSeq2Seq
from transformers import AdamWeightDecay
from transformers import AutoTokenizer, TFAutoModelForSeq2SeqLM
```

```
model_checkpoint = "Helsinki-NLP/opus-mt-en-hi"
```

```
# Loading the Dataset (Source: https://huggingface.co/datasets/cfilt/iitb-english-hindi)
raw_datasets = load_dataset("cfilt/iitb-english-hindi")

Repo card metadata block was not found. Setting CardData to empty.
WARNING:huggingface_hub.repocard:Repo card metadata block was not found. Setting CardData to empty.
```

```
# Dataset Info
raw_datasets
```

```
DatasetDict({
  train: Dataset({
    features: ['translation'],
    num_rows: 1659083
  })
  validation: Dataset({
    features: ['translation'],
    num_rows: 520
  })
  test: Dataset({
    features: ['translation'],
    num_rows: 2507
  })
})

raw_datasets['train'][1]

{'translation': {'en': 'Accerciser Accessibility Explorer',
                  'hi': 'एक्सेसइसर पहुँचनीयता अन्वेषक'}}

raw_datasets['test'][1]

{'translation': {'en': "As America's road planners struggle to find the cash to mend a crumbling highway system, many are beginning to see a solution in a little black box that fits neatly by the dashboard of your car.",
                  'hi': 'जबकि अमेरिका के सड़क योजनाकार, ध्वस्त होते हुए हाईवे सिस्टम को सुधारने के लिए धन की कमी से जूझ रहे हैं, वहीं बहुत-से लोग इसका समाधान छोटे से ब्लैक बॉक्स में देख रहे हैं, जो आपकी कार के डैशबोर्ड पर सफाई से फिट हो जाता है।'}}
```

▼ Preprocessing

```
# Initializing the Tokenizer
tokenizer = AutoTokenizer.from_pretrained(model_checkpoint)

/usr/local/lib/python3.10/dist-packages/transformers/models/mt5/tokenization_mt5.py:194: UserWarning: Recommended: pip install sac
warnings.warn("Recommended: pip install sacremoses.")

tokenizer(["I had about a 30 minute demo just using this new headset"])

{'input_ids': [[56, 154, 195, 19, 1671, 7336, 35914, 469, 1192, 90, 336, 1876, 8907, 0]], 'attention_mask': [[1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1]]}

with tokenizer.as_target_tokenizer():
    print(tokenizer(["मैं सिर्फ ३० minute का demo मि ला था इस नये headset का इस्तेमाल करने के लिए"])))

{'input_ids': [[4095, 14034, 3625, 44, 174, 18943, 3353, 14921, 6785, 3383, 39169, 24, 3947, 363, 818, 3245, 1754, 82, 89, 4075, 2326, 4
/usr/local/lib/python3.10/dist-packages/transformers/tokenization_utils_base.py:3660: UserWarning: `as_target_tokenizer` is deprecated &
warnings.warn(

# Tokenizing the English and Hindi words
max_input_length = 128
max_target_length = 128

source_lang = "en"
target_lang = "hi"

def preprocess_function(examples):
    inputs = [ex[source_lang] for ex in examples["translation"]]
    targets = [ex[target_lang] for ex in examples["translation"]]
    model_inputs = tokenizer(inputs, max_length=max_input_length, truncation=True)

    # Setup the tokenizer for targets
    with tokenizer.as_target_tokenizer():
        labels = tokenizer(targets, max_length=max_target_length, truncation=True)

    model_inputs["labels"] = labels["input_ids"]
    return model_inputs

preprocess_function(raw_datasets["train"][:1])

{'input_ids': [[3872, 85, 2501, 132, 15441, 36398, 0]], 'attention_mask': [[1, 1, 1, 1, 1, 1, 1]], 'labels': [[63, 2025, 18, 16155,
346, 20311, 24, 2279, 679, 0]]}

tokenized_datasets = raw_datasets.map(preprocess_function, batched=True)

model = TFAutoModelForSeq2SeqLM.from_pretrained(model_checkpoint)

All model checkpoint layers were used when initializing TFMarianMTModel.

All the layers of TFMarianMTModel were initialized from the model checkpoint at Helsinki-NLP/opus-mt-en-hi.
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFMarianMTModel for predictions with

# Initializing the Hyper Parameter
batch_size = 16
learning_rate = 2e-5
weight_decay = 0.01
num_train_epochs = 10

data_collator = DataCollatorForSeq2Seq(tokenizer, model=model, return_tensors="tf")
generation_data_collator = DataCollatorForSeq2Seq(tokenizer, model=model, return_tensors="tf", pad_to_multiple_of=128)

tokenized_datasets['train']
```

```

Dataset({
    features: ['translation', 'input_ids', 'attention_mask', 'labels'],
    num_rows: 1659083
})

# Splitting the Dataset to Train and Test data
train_dataset = model.prepare_tf_dataset(
    tokenized_datasets['test'],
    batch_size=batch_size,
    shuffle=True,
    collate_fn=data_collator,
)

validation_dataset = model.prepare_tf_dataset(
    tokenized_datasets["validation"],
    batch_size=batch_size,
    shuffle=False,
    collate_fn=data_collator,
)

generation_dataset = model.prepare_tf_dataset(
    tokenized_datasets["validation"],
    batch_size=8,
    shuffle=False,
    collate_fn=generation_data_collator,
)

# Declaring the optimizer for improving the accuracy and reduce the loss
optimizer = AdamWeightDecay(learning_rate=learning_rate, weight_decay_rate=weight_decay)
model.compile(optimizer=optimizer)

```

▼ Training the Model

```

model.fit(train_dataset, validation_data=validation_dataset, epochs=10)

Epoch 1/10
156/156 [=====] - 94s 375ms/step - loss: 3.7551 - val_loss: 3.9527
Epoch 2/10
156/156 [=====] - 49s 312ms/step - loss: 3.3230 - val_loss: 3.8702
Epoch 3/10
156/156 [=====] - 51s 325ms/step - loss: 3.0182 - val_loss: 3.8317
Epoch 4/10
156/156 [=====] - 50s 323ms/step - loss: 2.7760 - val_loss: 3.8207
Epoch 5/10
156/156 [=====] - 50s 320ms/step - loss: 2.5713 - val_loss: 3.8223
Epoch 6/10
156/156 [=====] - 49s 314ms/step - loss: 2.3807 - val_loss: 3.8225
Epoch 7/10
156/156 [=====] - 49s 316ms/step - loss: 2.2086 - val_loss: 3.8284
Epoch 8/10
156/156 [=====] - 50s 320ms/step - loss: 2.0565 - val_loss: 3.8544
Epoch 9/10
156/156 [=====] - 51s 328ms/step - loss: 1.9127 - val_loss: 3.8658
Epoch 10/10
156/156 [=====] - 51s 324ms/step - loss: 1.7789 - val_loss: 3.8789
<keras.callbacks.History at 0x7ee0130dabc0>

model.save_pretrained("tf_model/")

```

▼ Model Testing

```

from nltk.translate.bleu_score import sentence_bleu

tokenizer = AutoTokenizer.from_pretrained(model_checkpoint)
model = TFAutoModelForSeq2SeqLM.from_pretrained("tf_model/")

/usr/local/lib/python3.10/dist-packages/transformers/models/arian/tokenization_arian.py:194: UserWarning: Recommended: pip install sac
warnings.warn("Recommended: pip install sacremoses.")
All model checkpoint layers were used when initializing TFMarianMTModel.

```

All the layers of TFMarianMTModel were initialized from the model checkpoint at tf_model/.

If your task is similar to the task the model of the checkpoint was trained on, you can already use TFMarianMTModel for predictions with

```
raw_datasets['train']['translation'][1]['en']

'Accerciser Accessibility Explorer'

pred_input=raw_datasets['train']['translation'][1]['en']
tokenized = tokenizer([pred_input], return_tensors='np')
out = model.generate(**tokenized, max_length=128)
print(out)

tf.Tensor([[61949 26618 16155   346 33383    0 61949]], shape=(1, 7), dtype=int32)

with tokenizer.as_target_tokenizer():
    prediction=tokenizer.decode(out[0], skip_special_tokens=True)

# Input value
raw_datasets['train']['translation'][1]['en']

'Accerciser Accessibility Explorer'

# Original values
raw_datasets['train']['translation'][1]['hi']

'एक्सेसिबिलिटी एक्सप्लोरर'

# Predicted Value
prediction

'एक्सेसिबिलिटी एक्सप्लोरर'
```

▼ Testing with New Inputs

```
# Input 1
input_text = "Definitely share your feedback in the comment section."

tokenized = tokenizer([input_text], return_tensors='np')
out = model.generate(**tokenized, max_length=128)
with tokenizer.as_target_tokenizer():
    print(tokenizer.decode(out[0], skip_special_tokens=True))

टिप्पणी खंड में निश्चित रूप से अपनी प्रतिक्रिया साझा करें।
/usr/local/lib/python3.10/dist-packages/transformers/tokenization_utils_base.py:3660: UserWarning: `as_target_tokenizer` is deprecated &
warnings.warn(

# Input 2
input_text = "So even if it's a big video, I will clearly mention all the products."

tokenized = tokenizer([input_text], return_tensors='np')
out = model.generate(**tokenized, max_length=128)
with tokenizer.as_target_tokenizer():
    print(tokenizer.decode(out[0], skip_special_tokens=True))

तो यह एक बड़ा वीडियो है, तो भी मैं सभी उत्पादों का स्पष्ट रूप से उल्लेख करूँगा।

# Input 3
input_text = "I was waiting for my bag."

tokenized = tokenizer([input_text], return_tensors='np')
out = model.generate(**tokenized, max_length=128)
with tokenizer.as_target_tokenizer():
    print(tokenizer.decode(out[0], skip_special_tokens=True))

मैं अपने बैग के लिए प्रतीक्षा कर रहा था।
```

✓ 5s completed at 11:49 AM

● ×